

Loading in data

In [1]:

```
import numpy as np

from tensorflow.keras.utils import image_dataset_from_directory

from PIL import Image
```

In [23]:

```
train_ds = image_dataset_from_directory(
    "train/train",
    labels="inferred",
    label_mode="categorical",
    class_names=None,
    color_mode="rgb",
    batch_size=100,
    image_size=(500, 500),
    shuffle=True,
    seed=1,
    validation_split=.2,
    subset="training",
    interpolation="bicubic",
    follow_links=False,
    crop_to_aspect_ratio=False,
)

validation_ds = image_dataset_from_directory(
    "train/train",
    labels="inferred",
    label_mode="categorical",
    class_names=None,
    color_mode="rgb",
    batch_size=100,
    image_size=(500, 500),
    shuffle=True,
    seed=1,
    validation_split=.2,
    subset="validation",
    interpolation="bicubic",
    follow_links=False,
    crop_to_aspect_ratio=False,
)
```

Found 5656 files belonging to 5 classes.
Using 4525 files for training.
Found 5656 files belonging to 5 classes.
Using 1131 files for validation.

In [3]:

```
from sklearn.model_selection import train_test_split

from tensorflow.keras import datasets
from tensorflow.keras.utils import to_categorical
from tensorflow.keras.models import Sequential

from tensorflow.keras import layers

from tensorflow.keras.layers import Dense # creates densely connected layer object
from tensorflow.keras.layers import Flatten # takes 2D input and turns into 1D array

from tensorflow.keras.layers import Conv2D # convolution layer
from tensorflow.keras.layers import MaxPooling2D # max pooling layer
```

```
from keras.utils.vis_utils import plot_model
```

In [4]:

```
type(train_ds)
```

Out[4]:

tensorflow.python.data.ops.dataset_ops.BatchDataset

Base Model

In [5]:

```
rescaling_layer = layers.experimental.preprocessing.Rescaling(
    scale=1. / 255,
    input_shape=(500, 500, 3)
)
```

In [6]:

```
model = Sequential()

model.add(rescaling_layer)

model.add(Flatten())
model.add(Dense(64, activation="relu"))
model.add(Dense(5, activation="softmax"))

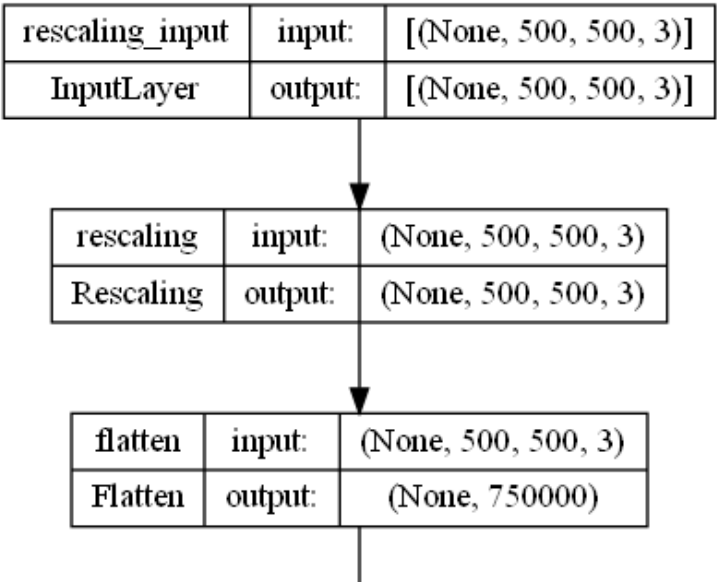
model.summary()
plot_model(model, show_shapes=True, show_layer_names=True)
```

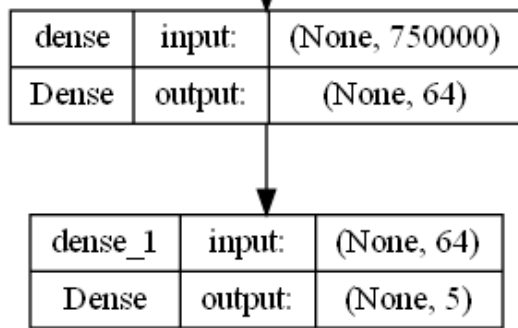
Model: "sequential"

Layer (type)	Output Shape	Param #
rescaling (Rescaling)	(None, 500, 500, 3)	0
flatten (Flatten)	(None, 750000)	0
dense (Dense)	(None, 64)	48000064
dense_1 (Dense)	(None, 5)	325

=====
Total params: 48,000,389
Trainable params: 48,000,389
Non-trainable params: 0

Out[6]:





In [8]:

```
model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

In [20]:

```
history_cnn = model.fit(train_ds, validation_data=validation_ds, epochs= 40,
                        validation_split = 0.2)
```

```
Epoch 1/40
71/71 [=====] - 21s 293ms/step - loss: 1.4068 - accuracy: 0.4707
- val_loss: 1.4098 - val_accuracy: 0.4668
Epoch 2/40
71/71 [=====] - 25s 340ms/step - loss: 1.3965 - accuracy: 0.4707
- val_loss: 1.4011 - val_accuracy: 0.4668
Epoch 3/40
71/71 [=====] - 25s 334ms/step - loss: 1.3877 - accuracy: 0.4707
- val_loss: 1.3934 - val_accuracy: 0.4668
Epoch 4/40
71/71 [=====] - 25s 331ms/step - loss: 1.3801 - accuracy: 0.4707
- val_loss: 1.3869 - val_accuracy: 0.4668
Epoch 5/40
71/71 [=====] - 25s 335ms/step - loss: 1.3736 - accuracy: 0.4707
- val_loss: 1.3815 - val_accuracy: 0.4668
Epoch 6/40
71/71 [=====] - 25s 339ms/step - loss: 1.3680 - accuracy: 0.4707
- val_loss: 1.3770 - val_accuracy: 0.4668
Epoch 7/40
71/71 [=====] - 29s 391ms/step - loss: 1.3633 - accuracy: 0.4707
- val_loss: 1.3731 - val_accuracy: 0.4668
Epoch 8/40
71/71 [=====] - 27s 366ms/step - loss: 1.3593 - accuracy: 0.4707
- val_loss: 1.3699 - val_accuracy: 0.4668
Epoch 9/40
71/71 [=====] - 26s 353ms/step - loss: 1.3559 - accuracy: 0.4707
- val_loss: 1.3672 - val_accuracy: 0.4668
Epoch 10/40
71/71 [=====] - 27s 361ms/step - loss: 1.3531 - accuracy: 0.4707
- val_loss: 1.3650 - val_accuracy: 0.4668
Epoch 11/40
71/71 [=====] - 26s 356ms/step - loss: 1.3507 - accuracy: 0.4707
- val_loss: 1.3632 - val_accuracy: 0.4668
Epoch 12/40
71/71 [=====] - 27s 362ms/step - loss: 1.3486 - accuracy: 0.4707
- val_loss: 1.3617 - val_accuracy: 0.4668
Epoch 13/40
71/71 [=====] - 29s 391ms/step - loss: 1.3469 - accuracy: 0.4707
- val_loss: 1.3605 - val_accuracy: 0.4668
Epoch 14/40
71/71 [=====] - 28s 376ms/step - loss: 1.3455 - accuracy: 0.4707
- val_loss: 1.3596 - val_accuracy: 0.4668
Epoch 15/40
71/71 [=====] - 27s 364ms/step - loss: 1.3444 - accuracy: 0.4707
- val_loss: 1.3588 - val_accuracy: 0.4668
Epoch 16/40
24/71 [=====>.....] - ETA: 15s - loss: 1.3332 - accuracy: 0.4753
```

KeyboardInterrupt

Traceback (most recent call last)

```
Input In [20], in <cell line: 1>()
----> 1 history_cnn = model.fit(train_ds, validation_data=validation_ds, epochs= 40,
  2 validation_split = 0.2)
```

File ~\anaconda3\envs\Mercury\lib\site-packages\keras\utils\traceback_utils.py:64, in filter_traceback.<locals>.error_handler(*args, **kwargs)

```
62 filtered_tb = None
63 try:
--> 64     return fn(*args, **kwargs)
65 except Exception as e: # pylint: disable=broad-except
66     filtered_tb = _process_traceback_frames(e.__traceback__)
```

File ~\anaconda3\envs\Mercury\lib\site-packages\keras\engine\training.py:1409, in Model.fit(self, x, y, batch_size, epochs, verbose, callbacks, validation_split, validation_data, shuffle, class_weight, sample_weight, initial_epoch, steps_per_epoch, validation_steps, validation_batch_size, validation_freq, max_queue_size, workers, use_multiprocessing)

```
1402 with tf.profiler.experimental.Trace(
1403     'train',
1404     epoch_num=epoch,
1405     step_num=step,
1406     batch_size=batch_size,
1407     _r=1):
1408     callbacks.on_train_batch_begin(step)
-> 1409     tmp_logs = self.train_function(iterator)
1410     if data_handler.should_sync:
1411         context.async_wait()
```

File ~\anaconda3\envs\Mercury\lib\site-packages\tensorflow\python\util\traceback_utils.py:150, in filter_traceback.<locals>.error_handler(*args, **kwargs)

```
148 filtered_tb = None
149 try:
--> 150     return fn(*args, **kwargs)
151 except Exception as e:
152     filtered_tb = _process_traceback_frames(e.__traceback__)
```

File ~\anaconda3\envs\Mercury\lib\site-packages\tensorflow\python\eager\def_function.py:915, in Function.__call__(self, *args, **kws)

```
912 compiler = "xla" if self._jit_compile else "nonXla"
914 with OptionalXlaContext(self._jit_compile):
--> 915     result = self._call(*args, **kws)
917 new_tracing_count = self.experimental_get_tracing_count()
918 without_tracing = (tracing_count == new_tracing_count)
```

File ~\anaconda3\envs\Mercury\lib\site-packages\tensorflow\python\eager\def_function.py:947, in Function.call(self, *args, **kws)

```
944 self._lock.release()
945 # In this case we have created variables on the first call, so we run the
946 # defunned version which is guaranteed to never create variables.
--> 947 return self._stateless_fn(*args, **kws) # pylint: disable=not-callable
948 elif self._stateful_fn is not None:
949     # Release the lock early so that multiple threads can perform the call
950     # in parallel.
951     self._lock.release()
```

File ~\anaconda3\envs\Mercury\lib\site-packages\tensorflow\python\eager\function.py:2453, in Function.__call__(self, *args, **kwargs)

```
2450 with self._lock:
2451     (graph_function,
2452      filtered_flat_args) = self._maybe_define_function(args, kwargs)
-> 2453 return graph_function._call_flat(
2454     filtered_flat_args, captured_inputs=graph_function.captured_inputs)
```

File ~\anaconda3\envs\Mercury\lib\site-packages\tensorflow\python\eager\function.py:1860, in ConcreteFunction.call_flat(self, args, captured_inputs, cancellation_manager)

```
1856 possible_gradient_type = gradients_util.PossibleTapeGradientTypes(args)
1857 if (possible_gradient_type == gradients_util.POSSIBLE_GRADIENT_TYPES_NONE
1858     and executing_eagerly):
1859     # No tape is watching; skip to running the function.
-> 1860 return self._build_call_outputs(self._inference_function.call(
1861     ctx, args, cancellation_manager=cancellation_manager))
1862 forward_backward = self._select_forward_and_backward_functions(
1863     args,
1864     ...
```

```

1864     possible_gradient_type,
1865     executing_eagerly)
1866 forward_function, args_with_tangents = forward_backward.forward()

```

```

File ~\anaconda3\envs\Mercury\lib\site-packages\tensorflow\python\eager\function.py:497,
in _EagerDefinedFunction.call(self, ctx, args, cancellation_manager)
    495 with _InterpolateFunctionError(self):
    496     if cancellation_manager is None:
--> 497         outputs = execute.execute(
    498             str(self.signature.name),
    499             num_outputs=self._num_outputs,
    500             inputs=args,
    501             attrs=attrs,
    502             ctx=ctx)
    503     else:
    504         outputs = execute.execute_with_cancellation(
    505             str(self.signature.name),
    506             num_outputs=self._num_outputs,
    (...)
    509             ctx=ctx,
    510             cancellation_manager=cancellation_manager)

```

```

File ~\anaconda3\envs\Mercury\lib\site-packages\tensorflow\python\eager\execute.py:54, in
quick_execute(op_name, num_outputs, inputs, attrs, ctx, name)
    52 try:
    53     ctx.ensure_initialized()
---> 54     tensors = pywrap_tfe.TFE_Py_Execute(ctx._handle, device_name, op_name,
    55                                         inputs, attrs, num_outputs)
    56 except core._NotOkStatusException as e:
    57     if name is not None:

```

KeyboardInterrupt:

Model 1

In [13]:

```

model1 = Sequential()

model1.add(rescaling_layer)

model1.add(Conv2D(
    filters= 32, kernel_size = (5, 5), activation = "relu",
    input_shape = (500, 500, 3)))

model1.add(MaxPooling2D(pool_size = (3, 3)))

model1.add(Conv2D(64, (5, 5), activation = "relu"))
model1.add(MaxPooling2D(pool_size = (3, 3)))
model1.add(Conv2D(64, (5, 5), activation = "relu"))

model1.add(Flatten())
model1.add(Dense(64, activation="relu"))
model1.add(Dense(5, activation="softmax"))

model1.summary()
plot_model(model1, show_shapes=True, show_layer_names=True)

```

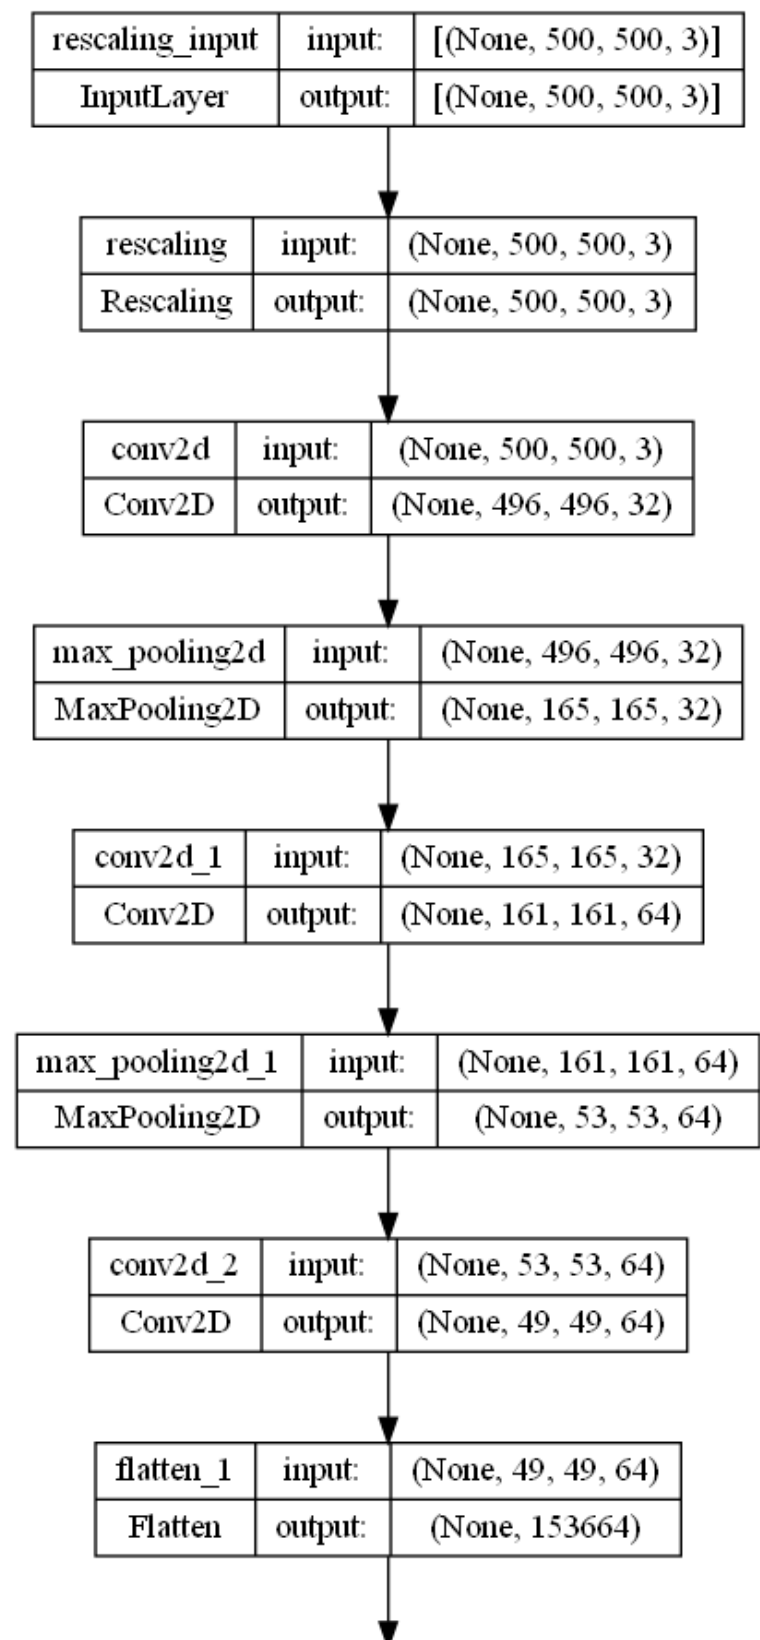
Model: "sequential_1"

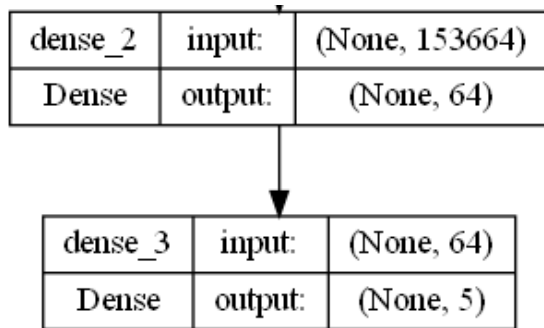
Layer (type)	Output Shape	Param #
=====		
rescaling (Rescaling)	(None, 500, 500, 3)	0
conv2d (Conv2D)	(None, 496, 496, 32)	2432
max_pooling2d (MaxPooling2D)	(None, 165, 165, 32)	0
)		
conv2d_1 (Conv2D)	(None, 161, 161, 64)	51264

max_pooling2d_1 (MaxPooling 2D)	(None, 53, 53, 64)	0
conv2d_2 (Conv2D)	(None, 49, 49, 64)	102464
flatten_1 (Flatten)	(None, 153664)	0
dense_2 (Dense)	(None, 64)	9834560
dense_3 (Dense)	(None, 5)	325

=====
Total params: 9,991,045
Trainable params: 9,991,045
Non-trainable params: 0

Out[13]:





In [21]:

```
model1.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

In [25]:

```
history_cnn = model1.fit(train_ds, validation_data= validation_ds, epochs= 1,
                        validation_split = 0.2)
```

```
46/46 [=====] - 710s 14s/step - loss: 0.0396 - accuracy: 0.9947
- val_loss: 3.6043 - val_accuracy: 0.5261
```

In []: